

Dudilig - AI PRD (Demo-Day Edition)

Product Name: Dudilig - The AI Copilot for Evaluating AI & Web3 Startups with Compliance Foresight

Author: Hersh Malik, Ken Ng

Date: October 10th 2025

What we learned from v1.0

1. **Too wide, too deep, too early.**

Our last PRD v1.0 with great research and a fully loaded roadmap at a Series-A company standard. We need to trim it down to a pre-demo prototype. V1.0 tried to serve VCs, CVCs, fintechs, and PE firms all at once, with seven frameworks and nine features. That will create diffusion and for the sake of our bootcamp demo, we would aim for more *validation*.

2. **The buyer is not ready to pay (yet).**

VCs want faster insight, not full compliance automation. They'll *admire* the concept but rarely budget for it. We should treat them as design partners, not as immediate revenue sources.

3. **Accuracy beats complexity.**

The single biggest risk in v1.0 was hallucination. For credibility, every output must have a source chip and confidence tag. Deterministic logic with human-readable citations matters more than fancy generative reasoning.

4. **"Predictive compliance" sells best when it touches deal terms.**

What gets attention isn't the risk score — it's showing how regulation changes valuation, timelines, and term sheets. That's the moment where a partner leans forward.

5. **QuantumLight showed the path.**

Revolut founder Nik Storonsky's QuantumLight proved that "AI-led diligence" can morph into fund-with-software. That's an optional future for us, but first we need one believable simulation that investors can see.

What we need to build for demo day (v1.2 MVP for Oct 17)

Goal: Ship a light, working prototype that ingests a pitch deck or URL and outputs a 1-page, citable **Compliance Impact Brief** in under **3 minutes**.

Scope (in Lovable/Replit terms):

Component	Description	Build type
Upload / Input	PDF or URL upload	Front-end React
Extractor	LLM-based NER for company, sector, geography, AI use-case	API (GPT-4 / Claude)
Classifier	Deterministic rules table mapping use-case → EU AI Act tier	Local YAML / CSV
Gaps View	4–5 checklist items with red/yellow/green badges + citations	Static data + Markdown render
Scenario (CIS Lite)	One preset “Stricter EU AI Act 2026” with cost/timeline/valuation deltas	JSON parameter sheet
Output	1-page PDF (summary, citations, term suggestions)	Server-rendered HTML→PDF
Style	Clean enterprise UI (white + slate)	TailwindCSS
Backup	Pre-baked demo dataset if API fails	In-memory fallback

Out of scope: multi-framework coverage, alerts, storage, auth, portfolio views, expert workflow.

Metrics:

- Time-to-brief < 3 min.
- 100 % of claims have citations.
- 0 demo crashes.
- 2–3 judges or advisors say “this would change how we evaluate a deal.”

0) One-liner

Turn a startup pitch deck into a cited EU AI Act risk brief and a single “what-if” scenario that changes deal terms — in under 3 minutes.

1) Purpose & Goal

- **Purpose:** Validate that VCs will use (and trust) a **cited, IC-ready** compliance brief that surfaces deal-term implications in minutes.
- **Demo-day goal (Oct 17):** A **working prototype** that ingests a pitch deck/URL and outputs:
 1. **Likely EU AI Act classification** with 2–3 citations,
 2. **4–5 concrete gaps** (R/Y/G) with citations,
 3. **One preset scenario** (“Stricter EU AI Act enforcement, 2026”) with **cost/timeline ranges** and **deal-term suggestions**,
 4. **Downloadable 1-page IC brief (PDF).**

Success signal: 5+ design-partner intros; at least 2 partners say the brief would **change valuation/terms/pass-no-pass**.

2) Target users & buyers (now → later)

- **Now (design partners):** VC **associates/senior associates** doing AI/health-data deals; need a fast, citable brief before partner review.
- **Later (economic buyer):** **Fintech/AI compliance teams** (CCO, Head of Risk) using the same simulator internally (larger budgets).
- **Optionality:** If we pivot toward **fund-with-software**, this same engine becomes our sourcing edge.

3) Problem (outcome framing)

Associates cannot translate regulation into **deal terms in <10 minutes**. They google, ping counsel, and still can’t answer:

- “What if enforcement tightens next year?”
- “How should we structure the round?”

We deliver a 1-pager that is immediately IC-usable and citable.

4) Product scope (Demo-day MVP)

In scope (must-have)

- **Input:** Deck upload (.pdf) or URL.
- **Extraction:** Company name, sector, geos, AI use case (LLM NER).
- **Classifier:** Likely EU AI Act tier (e.g., high-risk) via **deterministic rules table** (use-case × sector × geo).
- **Gaps:** 4–5 checklist items (e.g., CE marking plan, DPIA, bias eval, post-market monitoring) each with:
 - **Badge** (R/Y/G)
 - **Plain-English action**
 - **Source chip** (article/recital link)
- **CIS (single preset scenario): “Stricter EU AI Act enforcement 2026”**
 - **Cost range** (+\$450–650k/yr), **timeline range** (+8–12 mo), **deal-term suggestions** (reserve, milestone).
 - Each assumption has a tooltip and a source or industry reference.
- **Export:** 1-page “**Compliance Impact Brief**” (PDF).

Out of scope (cut for demo)

- Multiple regulations (GDPR/SEC/HIPAA), alerts, portfolio view, CRM/SSO, storage/auth, expert validation workflow, multi-scenario editor.

5) Value proposition (now)

- **Speed:** IC-ready brief in **<3 minutes**.
 - **Trust:** Every assertion has a **source chip**; scenario outputs are **ranges** with **confidence tags**.
 - **Impact:** Produces **deal-term changes** (reserve/milestone/valuation delta) not just red flags.
-

6) Core user journey (demo flow)

1. Drop a **curated healthcare-AI deck** (primary) or paste URL.
 2. See **Likely classification** (“High-Risk”) → tap to view **3 citations**.
 3. Review **4–5 gaps** with actions & sources.
 4. Click “**Simulate 2026 Stricter Enforcement**” → see **cost/timeline ranges + 2 term suggestions**.
 5. **Download 1-pager**; show how it plugs into IC memo.
-

7) Features (MVP spec)

F1. Extraction (LLM-assisted NER)

- Pull: company, sector (health AI), markets (US/EU), AI modality (CV/NLP), product claims.
- If extractor fails, fallback form (manual entry).

F2. Tier classification (deterministic)

- CSV/YAML **rules table** mapping common health-AI use cases to **likely EU AI Act tier**.
- Output includes **why** (+2–3 clause citations).

F3. Gap checklist (deterministic + templated)

- 10 canonical checks; show top 4–5 based on extracted context.
- For each: **status guess**, **action sentence**, **source chip**.

F4. CIS – single preset

- Fixed parameter sheet (staffing, audits, notified body, CE marking timelines) → **cost and delay ranges**.
- **Deal-term widget**: auto-compose 2 recommendations (reserve \$, milestone gate).

F5. PDF export

- Branded one-pager: classification, gaps, scenario deltas, recommendations, citations, date.

Non-goals for demo: persistence/auth, multi-user, alerts, confidence calibration beyond simple High/Med/Low tag.

8) Architecture (demo)

- **Front-end:** Single React page (upload → results).
 - **Back-end:**
 - **LLM** (NER + paraphrase only; no free-form legal claims).
 - **RAG** over **curated EU AI Act corpus** (official text + 1 explainer) for citations.
 - **Rules engine:** in-memory CSV/YAML (tiering, gaps, CIS params).
 - **PDF service:** server render.
 - **No DB/auth** (in-memory for demo).
 - **Panic toggle:** return pre-baked outputs if API stalls.
-

9) Data sources (demo)

- **EU AI Act:** official consolidated text (snippets with anchors).
- **Explainers:** 1–2 reputable public summaries (for secondary citations).
- **Assumption sheet:** salary/audit/CE costs; timeline ranges; each cell annotated with a reference.

10) Metrics (demo & pilot)

- **Time-to-brief:** <3 min from upload to PDF.
- **Decision influence (qual):** we record if a judge/advisor says it would **change valuation/terms/pass/no-pass**.
- **Citation coverage:** 100% of regulatory statements have a source chip.

- **Crash-free demo:** 100%.
-

11) Risks & mitigations (demo)

- **Hallucination:** Restrict LLM to NER; all regulatory claims come from **rules + citations**.
 - **Credibility:** Use **ranges** + “**Not legal advice**” banner + confidence tags.
 - **Flaky inputs:** Provide a **curated primary deck** and a **fallback deck**.
-

12) Timeline to Oct 17 (D-7 → D-0)

- **D-7:** Freeze scope; curate corpus; finalize **rules table** (use-cases × tier × gaps).
- **D-6:** NER extractor; fallback form.
- **D-5:** Classifier + citations view.
- **D-4:** Gap list with actions & chips.
- **D-3:** CIS card (ranges + term widget) wired to params.
- **D-2:** PDF export; confidence tags; styling pass.
- **D-1:** Rehearsal, screen-recorded backup, panic toggle.

13) Positioning

- **Now:** “**IC-ready compliance brief in minutes** for AI deals” (design partner with VCs).
- **Near-term option:** same simulator sold to **fintech/AI compliance teams** (real budgets).
- **Long-term option:** **fund-with-software** where Dudilig is the selection edge.

Tagline: “**From deck to deal terms — with citations.**”

14) Post-demo roadmap (short)

- **Pilot (Q1'26):** Add **GDPR + SEC crypto** coverage; 2 preset scenarios; PDF polish.
 - **Early fintech wedge:** templated scenarios for **model risk / KYC-AML** decisions.
 - **Trust:** light expert-review loop for critical flags; store analyses; simple auth.
-

15) What to show on stage (script cues)

- “Upload deck” → spinner (2–3 sec).
- “High-Risk under EU AI Act” → **tap citations**.
- “Top gaps” → one click per badge (plain English).
- “Simulate 2026” → **cost/timeline ranges** + “Add \$800k reserve; gate CE Marking pre-Series B.”
- “Export 1-pager” → flash the PDF preview.
- Close: “This saves **8–10 hours** and directly **changes term sheets**. VCs are our beachhead; the same simulator becomes a **decision copilot for fintech compliance**, where the real budgets live.”